Murmuration: Federating LLM Agents Across Machines with Nothing but Text

Discussion: questions and comments via GitHub Issues on this repository — no e-mail address is provided; open an issue instead.

Abstract

Personal LLM-agent harnesses now run simultaneously on many heterogeneous machines (workstations, cloud servers, LAN boxes, offline nodes), yet each instance is a silo: its sessions, files, and learned lessons do not cross machine boundaries, and letting one agent operate another machine's shell is a security red line. We present **Murmuration**, a method framework that federates *independent* agent instances through a single shared, plain-text context database, adding no broker or middleware beyond that store. Its core is a **semantic-level transport model**: messages carry only speech-act headers (five performatives), wire-level identities, and wall-clock timestamps; anything larger than a small threshold travels as a **content-addressed pointer** (path + SHA-256 + byte count) that is verified *outside* any agent's context window. We formalize a message calculus and derive three cost properties — design invariants implied by the payload threshold and the manifest discipline, not independently established empirical results: per-delivery context cost scales with the pointer list rather than payload bytes, and drops to a constant header when the manifest itself travels as a file; novelty detection is O(1) via fixed-width timestamp comparison; and k-hop relay never re-pays payload context, since pointer forwarding and digest verification are idempotent. Around the transport we specify organizational governance: a versioned, incident-driven handbook with red lines, role-scoped acknowledgment obligations, and time-triggered watchkeeping suited to turn-based agents. We report a sanitized eleven-day (elapsed) field deployment across five heterogeneous machines (protocol v1.0→v1.4, eight role types, 45 cross-machine lessons, 13 archived incidents, deliveries of up to 1,300 files, each verified by per-file SHA-256 corruption checks) and a 49-system novelty assessment. Within the window the room carried 283 structured messages (mean 26/day, peak 49/day) — 39 typed data deliveries, 27 asks, 87 acknowledgment discharges — with more than three thousand per-file SHA-256 checks at byte equality and zero mismatches; payload content (scientific-experiment pipelines, document drafting, paper writing) never enters this paper, only counts and shapes. Per our search (2026-11), no public scheme combines all of: heterogeneous cross-machine instances, a shared plain-text context database as the sole substrate, versioned incident-driven handbook governance, pointer-based delivery-integrity checks, turn-driven watchkeeping, and field evidence; we state claims at a deliberately conservative level.

Keywords: multi-agent systems, LLM agents, agent communication, semantic communication, content addressing, coordination artifacts, organizational governance, field deployment.

1. Introduction

A single developer or small team today routinely operates several LLM-agent instances [1]: one on a laptop, one on a cloud server, one on a LAN machine without external network, and so on. Each instance is *turn-based* (it acts, returns, and is re-invoked), *context-bounded* (its working memory is a token window that is billed and that degrades with length), and *environment-privileged* (it holds credentials and shell rights on exactly one machine). Three consequences follow.

1. **Silos.** Lessons learned on machine A (a command trap, a config pitfall, an incident root cause) are invisible to machine B unless a human copies them.

2. **Context scarcity as the true bandwidth.** When agents collaborate, the scarce channel is not network bandwidth but the context window [2,3,4]: every byte that enters it is paid for in tokens and competes with the task at hand.
3. **The remote-shell red line.** The naive fix — letting agent A ssh into machine B and run commands — merges trust domains and is exactly what operational red lines forbid ("never operate another machine on its behalf").

Existing answers address *single-host* multi-agent work (built-in agent teams, swarm frameworks, conversation orchestrators) [5,6,7,8,9] or *platform-scale* infrastructure and cross-machine agent protocols (organization-level substrates with identity, discovery, and container orchestration) [10,11]. Neither fits the personal/small-team setting: independent heterogeneous instances, no open ports, no broker clusters, no budget for platform operations, and a human who must be able to audit everything by reading plain text.

Positioning as a research prototype. Murmuration is a research prototype for cross-architecture multi-agent collaboration and work relay (cross-harness by design): independent agent instances on different machine architectures — and different builds of the DeepSeek Harness (DSH) family [9] — hand work to one another through the federation, and compared with traditional single-session execution it delivers the same class of multi-agent coordination at near-zero marginal infrastructure cost in our deployment (no broker, no open ports, no platform team). Concurrently, the DeepSeek Harness team's Super-agent design exploration [12] reports the same attention bottleneck — a human sustains only 3–5 active sessions — and answers it inside one organization's footprint with a two-layer Router/Worker architecture, event-driven triggers with thread subscription, Inline & Spill and Chain-of-Recall context mechanics, speculative execution, and a Watchdog control loop, while listing cross-instance agent-to-agent (A2A [10]) communication and cross-session human collaboration as open problems. We occupy exactly that open slot from the opposite direction: rather than redesigning the interior of one agent stack, we federate already-deployed independent instances with plain-text organizational artifacts — nothing but text in the coordination plane, since payload bytes move over pre-existing point-to-point channels — across machines and harness builds.

Contributions. (i) A method framework that federates independent, heterogeneous LLM-agent instances over one shared plain-text context database, adding **no broker or middleware beyond that store** (§3). (ii) A **semantic-level transport model** with a message calculus and three cost propositions (§4), positioned against speech-act communication, tuple spaces, publish/subscribe with consumer offsets, content addressing, and provenance standards. (iii) An **organizational governance layer**: versioned incident-driven handbook, red-line evolution, role-scoped acknowledgment obligations, and time-triggered watchkeeping for turn-based agents (§5). (iv) A **sanitized field study** — a real deployment instrumented with sanitized counts (five machines, eleven days) — plus a 49-system × 10-dimension novelty assessment with an explicit claim ladder (§6, §7).

2. Setting and Constraints

Agents. Each node is a turn-based LLM-agent harness. All five field nodes run DeepSeek Harness (DSH) builds [9]; the protocol requires only UTF-8 file I/O and is designed to be harness-agnostic, although a second harness family has not yet been fielded. An agent cannot block: between turns it is inert. Hence all liveness must come from *scheduled re-invocation*, not from listening sockets.

Substrate. One shared context database reachable by every node through *some* path: an outbound ssh tunnel, a local/sshfs mount, or — for a node with no network at all — a human or peer relay

carrying files. The substrate stores UTF-8 text; binary payloads live in an attached store and are referenced, never embedded.

Threat and audit model. No inbound ports; no agent ever executes commands on another machine; every coordination log — the public channel, the machine inboxes, and the role cursors — is append-only plain text that a human can read end-to-end (read-only governance texts such as guides are maintained by version-numbered replacement rather than append); integrity rests on SHA-256 digests computed outside agent contexts (a corruption check, not an authenticity control; see §8). Credentials never appear in the shared substrate.

Heterogeneity. Nodes differ in OS, model backend, network reachability, and uptime. The only required capability is reading and writing UTF-8 text files.

3. The Room: Coordination Artifact as Transport

The federation's coordination state lives in a *room*: a directory of append-only Markdown logs.

- `general.md` — the public channel;
- `inbox/<machine>.md` — per-machine inbound queue;
- `cursor/<machine>-<role>.md` — per-role consumption offsets;
- role and governance texts elsewhere in the shared database.

A **message** is a fenced block with a YAML-like header and a body:

```
---
mid: <machine>-<role>-<YYYYMMDDTHHMMSS>
from: <machine>
role: <role>
to: <machine> | general
ts: <ISO-8601>
type: ask | reply | note | data | ack
re: <mid>            # mandatory for reply
---
<body, or pointer lines for type=data>
```

Field authority. `mid` is the sole authority for identity, ordering, and deduplication; the human-readable `ts` header is advisory metadata for human audit and may drift under clock skew — consumers never compare `ts` values. The uniqueness scope of `mid` is (producer, second): two messages from one (machine, role) within the same second would collide — a hazard acknowledged in §8(iv).

Design rules from the deployment history: R1, R2, R3, and R5 are direct patches for recorded incidents, while R4 and R6 are stipulated design constraints (§6):

- **R1 (append-only).** Logs are never rewritten; corrections are new messages. Provenance is structural.
- **R2 (seconds-granular ids).** `mid` embeds a fixed-width timestamp *with seconds*; two roles on one machine posting in the same minute collided before this rule.
- **R3 (tail-timestamp novelty).** "Have I seen this?" compares only the fixed-width timestamp suffix of `mid`. Comparing whole ids lexicographically is a trap: machine-name prefixes reorder ids across machines.

- **R4 (payload threshold).** Bodies above 2 KB must become `data` messages whose body is pointer lines `data: <path> sha256=<hex> bytes=<n>`.
- **R5 (one cursor per role).** Cursors are per (machine, role); sharing one cursor across roles on a multi-role machine lost messages. The single cursor spans all channels a role scans; the cross-channel skip hazard is acknowledged in §8.
- **R6 (no secrets).** The substrate carries no tokens, keys, or personal data; payloads with such content stay out of the shared store entirely.

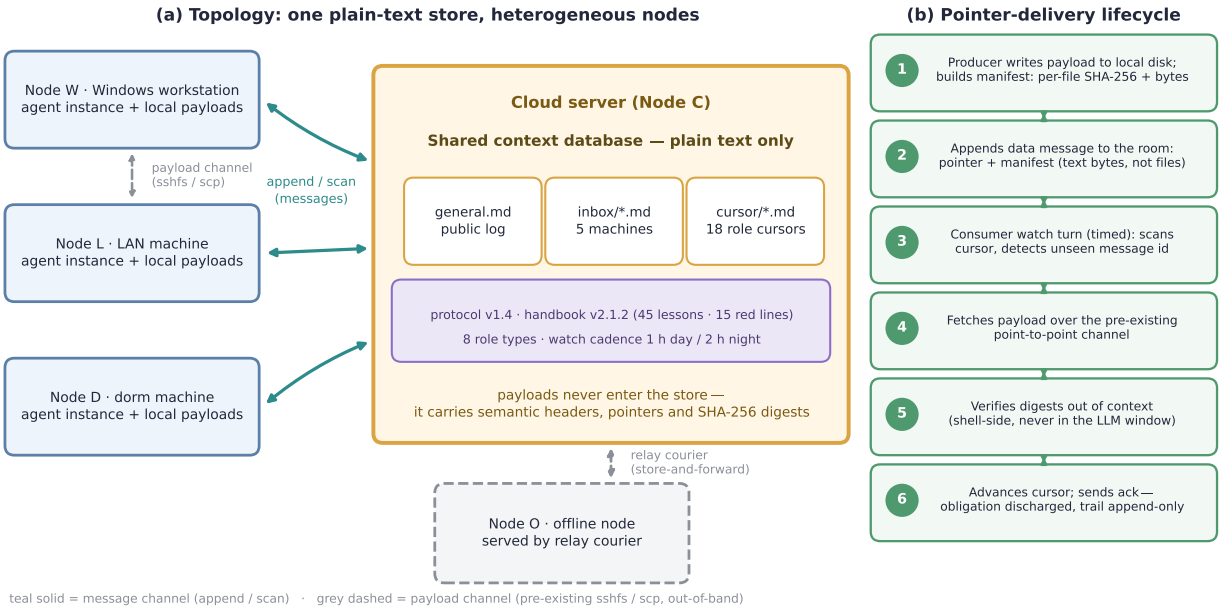


Figure 1: **Figure 1.** Murmuration at a glance. (a) Heterogeneous nodes coordinate through one shared plain-text store — public log, per-machine inboxes, per-role cursors — under a versioned governance handbook and time-triggered watchkeeping; payloads never enter the store and move over pre-existing point-to-point channels. (b) Pointer-delivery lifecycle: manifests and digests travel as text inside the context budget; files are fetched out-of-band and verified shell-side.

4. Semantic-Level Transport Model

4.1 Message calculus

Definition 1 (semantic message). A message is a triple $d = (a, p, r)$ where a is the *act header* (performative $\pi \in \{\text{ask, reply, note, data, ack}\}$, identifier *mid*, unique per (producer, second), identity $(m_{from}, \rho_{from}, m_{to})$ — the sending machine and posting role; the target role is resolved at delivery time from the addressed inbox and is not carried on the wire — a per-producer wall-clock timestamp t with bounded (NTP-grade) skew, and an optional referent re), p is an optional *content-addressed pointer* (path, h, n) with $h = \text{SHA-256}$ and byte count n, and r is an optional inline residue with $|r| \leq B$ (deployment value $B = 2$ KB).

Definition 2 (channel). A channel C is an append-only sequence $d_1, d_2, \ldots$; its *semantic projection* is the sequence of headers and pointers only.

Definition 3 (cursor). For a consumer (m, ρ), the cursor $\kappa(m, \rho)$ records the identifier — and hence the embedded fixed-width timestamp — of its last consumed message; novelty is decided by

R3's suffix comparison against that timestamp. The cursor is *per consumer* (one file per machine–role pair) and is compared against every channel the consumer scans (its inbox and the public channel); it remains well-defined under partial replication (an offline node's replica). One shared last-consumed identifier across two channels can skip older messages in the second channel — the cross-channel skip hazard of §8 — so consumers scan in a fixed channel order and accept a residual window.

Definition 4 (delivery manifest). A delivery of payload set $P = \{f_1, \ldots, f_k\}$ is one `data` message whose pointer list enumerates $(\text{path}_i, h_i, n_i)$; acceptance requires, for each i, an out-of-context equality check SHA-256$(f_i) = h_i \wedge |f_i| = n_i$, after which the receiver posts `ack`.

Definition 5 (watch schedule). A function $W : \text{time} \to \{\text{invoke}, \text{skip}\}$ re-invokes each role on a cadence (deployment: 1 h daytime, 2 h night), implementing time-triggered activation for agents that cannot listen.

4.2 Cost model

Let H be header+pointer size (hundreds of bytes) and $|P|$ payload size.

Proposition 1 (context cost). Assume manifests are generated and digest-verified entirely out of context, with only per-file boolean/summary lines entering it (the *manifest assumption*; this is exactly how the deployment's verification scripts run). Then the context-window cost charged to sender and receiver for one delivery of k files is $O(k \cdot H)$, independent of payload bytes $|P|$; and when the pointer list itself exceeds the inline threshold it is delivered as a file — one pointer, $O(H)$ — restoring a constant per-delivery context cost at any k. *Sketch.* By R4 the payload never enters any context; only a and p do, plus the verifier's summary. A manifest that is itself a payload is subject to R4 again, which yields the two-case bound. ■

Proposition 2 (O(1) novelty). Given cursor κ, and assuming no two messages from the same producer share a timestamp second (the uniqueness scope of *mid*; §8(iv)), detecting whether new messages exist costs one fixed-width string comparison per scanned channel, not a scan of message bodies. *Sketch.* Under that uniqueness assumption, R3's suffix is a total order key; bodies are irrelevant to novelty. ■

Proposition 3 (relay idempotence). Forwarding a pointer across k hops charges $O(H)$ context per hop, and the payload context is paid zero times beyond the original posting: the pointer denotes the same content at every hop, and verification remains end-to-end by digest equality. *Sketch.* Content addressing makes p location-independent; under the standard assumption that SHA-256 is collision-resistant, h ties pointer identity to content at every hop. ■

These properties map directly onto the Multi-Agent System Failure Taxonomy (MAST) [13]: specification ambiguity is bounded by five performatives with mandatory referents; inter-agent misalignment is bounded by explicit ack obligations (§5); unrecoverable failures become visible because every act is logged append-only; and context poisoning/overflow is bounded by Proposition 1. We state this as *design intent*, not as measured efficacy; incident-level grounding (incident → patching rule) ships as a repository artifact (§6).

4.3 Lineage

Speech acts [14] → KQML [15] → FIPA-ACL [16] supply the performative discipline; Linda and blackboards [17,18] supply shared coordination artifacts; publish/subscribe with consumer offsets (Kafka logs) [19,20] supply cursors; Merkle trees and IPFS content addressing [21,22] and W3C

PROV [23] supply pointers and provenance; event- versus time-triggered architecture [24] supplies the watch schedule; goal-oriented semantic communication [25,26] supplies the "transmit intent, not bits" stance; and stigmergy [27] and coordination theory [28] frame the room itself as an indirect-coordination medium. Classical ACL protocols assume that receiving is cheap: content travels inline and receiver cost is the sender's concern. Murmuration inherits their performative discipline but inverts the cost model for context-bounded receivers — payload bytes are replaced by content-addressed pointers verified outside context, and the R4 threshold caps what any single header can charge a receiver's window. Which discipline is inherited and which constraint is new is exactly the distinction the L2 claim (§7) rests on. The combination — and its reduction to *plain files* as the universal narrow waist — is what we contribute.

5. Organizational Governance

Transport without governance decays: roles blur, obligations lapse, lessons evaporate. Murmuration specifies four instruments, in the tradition of electronic institutions and organizational design for MAS [29,30].

Roles. Named posts per machine (diplomat, mechanic, editor, coordinator, ...); a machine may hold several; every message header carries the posting role, making responsibility auditable.

Versioned handbook with red lines. Cross-machine lessons live in an append-only handbook with numbered entries and an index; operational prohibitions are *red lines* with monotonic numbering (never renumbered). Crucially, red lines are **incident-driven patches**: each of the 13 archived incidents from the deployment history (one predating the study window) produced at least one rule (e.g., a session-registry overwrite incident produced the "never rewrite coordination state" family; a tunnel-death incident produced the watchdog rule). The handbook is thus a living patch series over the organization's failure surface.

Acknowledgment obligations. Borrowing from commitment-based MAS semantics [31]: a `data` or `ask` message creates a debtor obligation on its addressee, discharged only by `ack` or `reply`. A derived **ack board** (two columns: inbound obligations uncleared; outbound messages awaiting discharge) is computable from the logs alone and drives the watch turns.

Time-triggered watchkeeping. Because agents cannot block (Definition 5), liveness is a schedule. Cadence is a policy parameter (day/night tiers), and each watch turn is a fixed checklist: scan cursors, discharge obligations, update the board, escalate staleness.

Recovery. All organizational state is rebuildable from the logs alone. A lost or corrupted cursor is repaired by rescanning its channel from the first message and re-deriving the last-consumed identifier; the duplicate re-dispatch this may cause is harmless because acknowledgment discharge is idempotent per referent (an `ack` naming an already-discharged `mid` is a no-op), so recovery costs one linear scan and never payload re-verification unless a digest is re-challenged. A restarted watcher resumes from its cursor file; if that file is missing too, it adopts the channel tail only after posting a human-visible `note` announcing the reset, so the operator can veto silent state loss.

6. Field Study (sanitized)

Five heterogeneous machines (a Windows workstation, a cloud Linux server hosting the shared database, a LAN node, a dormitory node, and an offline node served by relay) ran the method for eleven elapsed days (2026-10-22 → 11-02; twelve calendar days inclusive). Machine names are pseudonyms; all identifiers, paths, and contents are sanitized.

- **Protocol evolution:** v1.0 → v1.4 within the eleven-day window, all changes backward-compatible; each minor version traceable to incidents or onboarding friction.
- **Organization:** eight role types; multi-role machines; one editor-in-chief owns the shared handbook (v2.1.2: four volumes, 45 cross-machine lessons, 15 red lines, 13 incident archives).
- **Deliveries:** thousand-file-scale packages (chapter drafts, figure sets, CSV matrices) moved as pointer manifests with per-file SHA-256 + byte verification — a corruption check, not an authenticity control (§8); zero digest mismatches across all verified files; per-delivery context cost stayed at header scale in operator observation (not separately metered).
- **Volume (sanitized counts only):** 283 structured messages across one public channel, five machine inboxes and eighteen role cursors during the window (mean 26/day, peak 49/day); 39 `data` deliveries, 27 `ask`s, 87 `ack`/`reply` discharges, 121 status notes, and 9 governance posts (four notices, a request, an action marker, a report, a correction); more than three thousand per-file SHA-256 digests checked with byte equality in total, largest single manifest above 1,300 files.
- **Task classes (abstracted):** scientific-experiment data pipelines, long-form document drafting, and paper writing; no payload content is disclosed — only counts, shapes and verification outcomes.
- **Liveness:** watch cadence institutionalized (1 h day / 2 h night); a 48-hour escalation rule for unanswered obligations; weekly digest.
- **Failure pedagogy:** the room rules R2, R3, and R5 each exist because a real message was lost or misordered without them; a related message-fence parsing rule was adopted later (2026-11-08) from a local smoke-test regression outside the study window and never caused a room-level loss.

A natural ablation. Because every rule is a patch for a recorded incident, the deployment history is a natural ablation: each rule partitions the timeline into a before-period with an observed failure and an after-period without recurrence. In the window, the incident-derived room rules (R2, R3, R5) showed no recurrence after adoption, and no incident recurred after its ops red line (the machine-layer rule family, numbered separately from the room rules) was written; per-rule exposure denominators are not separately metered, so these are narrative, not causal, claims. The full counterfactual tables and honesty caveats ship as a repository artifact (`docs/research/natural-ablation.md`).

We report this as *existence evidence* for operability, not as a controlled evaluation.

Artifact availability. The framework source, plugin tools, and the research artifacts behind §7 — the 49-system × 10-dimension novelty matrix, the search log, the adversarial re-review, and the claim ladder — ship in the public repository accompanying this paper (github.com/xuanjat/Murmuration, under `docs/research/`); documentation is licensed CC BY 4.0.

7. Novelty Assessment and Claim Ladder

We compared 49 systems across 10 dimensions — (1) cross-machine/network topology, (2) coordination substrate, (3) longevity (persistent organization vs. single task), (4) role systems including multi-role machines, (5) governance mechanisms, (6) delivery-integrity protocols, (7) watchkeeping/wake-up models, (8) human-in-the-loop boundaries, (9) privacy-boundary discipline, and (10) real operational evidence — spanning single-host orchestrators, cross-machine protocols, Claude/DSH-ecosystem tools, memory systems, academic simulations, and governance toolkits; the full matrix, search log, and adversarial re-review ship as repository artifacts.

Nearest neighbors: (a) **Avernet** [11], a platform substrate for organization-scale persistent agents

(identity/discovery/routing/containers; since 2026-09 it also ships a DSH channel plugin) — a different route: the distinction is administrative footprint and auditability, not the presence of infrastructure — the shared store is itself a hosted database service, but one append-only text store is a far lighter administrative footprint than identity, discovery, routing, and container orchestration, and every coordination artifact remains human-readable; (b) **ruflo** (formerly claude-flow) [8], a single-box swarm framework with hierarchical memory — complementary, not cross-machine; (c) **built-in single-host agent teams** [9] and conversation-based orchestrators [5,6,7] — complementary; (d) **concurrent single-organization super-agent architecture** (two-layer Router/Worker, event triggers with thread subscription, Inline & Spill and Chain-of-Recall context mechanics, Watchdog control loop; prototype plugin `dsh-superagent-prototype`) [12] — same motivation (attention bottleneck, session silos), different unit of redesign: they re-engineer one organization's agent interior with programmatic routing and push events on a single harness deployment, whereas we leave every instance untouched and add only a cross-instance organizational layer that works across machines and harness builds; their stated open problems (A2A between agent instances, collaboration detached from sessions) are precisely the surfaces our room protocol, pointer delivery, and ack obligations specify. Cross-machine agent protocols such as A2A [10] standardize task exchange between reachable service endpoints but do not specify governance, delivery integrity, or watchkeeping for turn-based personal instances behind NAT or on offline links. (e) Per-element closest neighbors from the repository matrix: cross-machine topologies — Avernet [11], IoA, AgentScope, ANP, A2A (all presuppose service endpoints or online connectivity); shared-substrate coordination — Letta's shared memory blocks and the mcp-agent-mail shared inbox (single-host; no delivery-integrity protocol); governance — a policy/compliance governance toolkit and academic institutional design (no versioned, incident-driven handbook evolution); delivery integrity — governance-toolkit audit trails (closest neighbor, but a different purpose: retrospective compliance logging, not receiver-side acceptance); watchkeeping — night-shift batch schedulers and cron-style wake-ups (scripts, no persistent organization or ack obligations); field evidence — simulation studies such as Project Sid (not a deployment).

Per our search (2026-11), no public scheme combines heterogeneous cross-machine instances + a shared plain-text context database as the sole substrate (in this design) + versioned incident-driven handbook governance + pointer-based delivery-integrity protocol + turn-driven time-triggered watchkeeping + field evidence. Our claim ladder — a repository discipline with an explicit weakest-to-strongest ordering: L0 (unqualified "first-ever" claims) and L1 ("first open-source" claims) are forbidden as unverifiable by exhaustive search, L2 (qualified combination novelty) is the ceiling this paper uses, L3 (contribution statements without uniqueness) fills the contributions list, and L4 (reusability-only) is the retreat position under challenge — states results at level L2 ("combination unseen in search") with contribution-style L3 sentences; "first-ever" claims are explicitly disclaimed, and the strongest single dimension is delivery integrity (dimension 6: digest discipline outside context), the weakest are cross-machine topology (dimension 1: every surveyed scheme already federates machines in some form) and human-in-the-loop boundaries (dimension 8: standard interrupt/approval support).

8. Limitations

No push delivery (polling only); no server-side deduplication or permission isolation in the substrate (compensated by cursor discipline and red lines); a single shared trust root; parameters (2 KB, cadences) are deployment choices, not optima; metrics are self-reported from one organization; the MAST alignment is design intent, not measured efficacy; the comparison with concurrent super-agent

designs [12] is qualitative — their measurements run on synthetic loads, ours on sanitized field logs, and no head-to-head benchmark exists.

Distributed-systems candor. (i) *Concurrent appends:* R1's history guarantee assumes appends do not interleave; network filesystems generally provide neither atomic append nor locking. The deployment compensates with role discipline (per-machine files, diplomat-mediated posting to the public channel) — discipline, not a filesystem guarantee. The substrate's per-append atomicity is in fact what the safety argument leans on: inboxes remain writable by any machine, so acknowledgment fan-in from many receivers to one inbox is serialized only by that property, and §3's history guarantee is conditional on it. (ii) *Integrity versus authenticity:* SHA-256 digests detect corruption, not adversarial substitution; any writer with substrate access could replace a payload and recompute its digest. The threat model of §2 presumes a single, mutually trusted trust root. (iii) *Injection bus:* message bodies enter peers' contexts verbatim, so a compromised node could attempt to steer the others; current mitigations are procedural (red lines, human-auditable plain text), not a data/instruction separation mechanism — an open problem for LLM federations generally. (iv) *Clocks and collisions:* R3's ordering presumes bounded (NTP-grade) clock skew across machines, and seconds-granular `mid`s can still collide for two same-second messages from one (machine, role). (v) *Relay consistency:* replica merging and duplicate handling for offline nodes rest on cursor discipline, without a formal reconciliation protocol. (vi) *Unbounded growth:* append-only logs have no archival or garbage-collection policy; watch-turn scans grow with the log. (vii) *Liveness bound:* worst-case obligation latency is the watch cadence plus the 48-hour escalation window; nothing tighter is guaranteed. (viii) *Cursor races and channel skips:* a cursor file is read, compared, and rewritten by its owner, and two concurrent watch turns of the same role can race on that read-modify-write — mitigated by one-file-per-role and turn scheduling, not eliminated; the single cursor also spans both channels a role scans, so advancing it on one channel can silently skip older messages on the other (the cross-channel skip hazard of Definition 3), narrowed by fixed channel order. (ix) *Parsing fragility:* message boundary detection depends on fence handling; a substrate append that drops a trailing newline can glue the next header into the previous body — the 2026-11-08 local regression behind the message-fence parsing rule (§6), mitigated by blank-line padding and read-back self-checks. (x) *Offline clock discipline:* the relay-served offline node's clock is not NTP-disciplined, so its messages can be judged stale by R3 comparisons; the skew assumption of (iv) is unverified for that node.

9. Conclusion

Murmuration suggests that a federation of independent, heterogeneous LLM agents can be coordinated in the coordination plane with *nothing but text* — payload bytes move over pre-existing point-to-point channels (§2): semantic headers and content-addressed pointers keep the scarce resource — context — free of payloads, while a versioned handbook, ack obligations, and time-triggered watchkeeping keep the organization alive between turns. The eleven-day, five-machine deployment indicates that the design is not merely implementable on paper but operable in sustained use, and its incident-driven rule set is a reusable artifact for anyone running personal agent federations.

References

1. Guo, T., et al. (2024). Large language model based multi-agents: a survey of progress and challenges. *arXiv:2402.01680* (IJCAI 2024).
2. Zhang, Z., et al. (2024). A survey on the memory mechanism of large language model based agents. *arXiv:2404.13501.*

3. Liu, N. F., et al. (2024). Lost in the middle: how language models use long contexts. *TACL* (arXiv:2307.03172).
4. Chen, L., Zaharia, M., & Zou, J. (2023). FrugalGPT: how to use large language models while reducing cost and improving performance. *arXiv:2305.05176*.
5. Wu, Q., et al. (2023). AutoGen: enabling next-gen LLM applications via multi-agent conversation. *arXiv:2308.08155*.
6. Hong, S., et al. (2023). MetaGPT: meta programming for a multi-agent collaborative framework. *arXiv:2308.00352*.
7. Qian, C., et al. (2023). Communicative agents for software development. *arXiv:2307.07924*.
8. ruflo (formerly claude-flow) (2026). *Single-box agent swarm orchestration with hierarchical memory*. https://github.com/ruvnet/ruflo (repository, accessed 2026-11).
9. DeepSeek (2026). *DeepSeek Harness (DSH)* — a local-first LLM-agent harness with built-in single-host Agent Teams — and *OpenViking*, its shared context database. npm package `@deepseek-ai/dsh`: https://www.npmjs.com/package/@deepseek-ai/dsh (accessed 2026-11).
10. Linux Foundation (2025). *Agent2Agent (A2A) protocol specification*.
11. inclusionAI (2026). *Avernet: infrastructure for persistent, organization-scale multi-agent systems*. https://github.com/inclusionAI/Avernet (repository documentation, accessed 2026-11).
12. DeepSeek Harness Team (2026). *Super-agent design exploration: toward more autonomous agent systems* (in Chinese) (architecture exploration article, 2026-11-01; prototype: dsh-superagent-prototype). Coverage: DoNews, https://www.donews.com/news/detail/1/6760986.html; PChome, https://article.pchome.net/news/16524.html (both 2026-11-01).
13. Cemri, M., et al. (2025). Why do multi-agent LLM systems fail? *arXiv:2503.13657* (NeurIPS 2025 Datasets & Benchmarks).
14. Searle, J. R. (1969). *Speech Acts: An Essay in the Philosophy of Language*. Cambridge University Press.
15. Finin, T., Fritzson, R., McKay, D., & McEntire, R. (1994). KQML as an agent communication language. In *Proceedings of CIKM '94*. ACM.
16. FIPA (2002). *FIPA Communicative Act Library Specification*, document no. SC00037F. Foundation for Intelligent Physical Agents. http://www.fipa.org/specs/fipa00037/.
17. Gelernter, D. (1985). Generative communication in Linda. *ACM TOPLAS*, 7(1).
18. Nii, H. P. (1986). Blackboard systems. *AI Magazine*, 7(2).
19. Eugster, P. Th., Felber, P. A., Guerraoui, R., & Kermarrec, A.-M. (2003). The many faces of publish/subscribe. *ACM Computing Surveys*, 35(2).
20. Kreps, J., Narkhede, N., & Rao, J. (2011). Kafka: a distributed messaging system for log processing. *NetDB '11*.
21. Merkle, R. C. (1980). Protocols for public key cryptosystems. *IEEE Symposium on Security and Privacy*.
22. Benet, J. (2014). IPFS — content addressed, versioned, P2P file system. *arXiv:1407.3561*.
23. Moreau, L., et al. (2013). *PROV-Overview: An Overview of the PROV Family of Documents*. W3C.
24. Kopetz, H. (1991). Event-triggered versus time-triggered real-time systems. In *Operating Systems of the 90s and Beyond* (Dagstuhl Seminar), LNCS 563. Springer.
25. Shannon, C. E., & Weaver, W. (1949). *The Mathematical Theory of Communication*. University of Illinois Press.
26. Luo, J., Delfani, E., Salimnejad, M., & Pappas, N. (2025). From information freshness to semantics of information and goal-oriented communications. *arXiv:2512.12758*.
27. Theraulaz, G., & Bonabeau, E. (1999). A brief history of stigmergy. *Artificial Life*, 5(2).

28. Malone, T. W., & Crowston, K. (1994). The interdisciplinary study of coordination. *ACM Computing Surveys*, 26(1).
29. Esteva, M., Rodríguez-Aguilar, J. A., Sierra, C., García, P., & Arcos, J. L. (2001). On the formal specification of electronic institutions. *AMEC III*, Springer.
30. Hübner, J. F., Sichman, J. S., & Boissier, O. (2007). Moise+: towards a structural, functional, and deontic model for multi-agent organizations. *International Journal of Agent-Oriented Software Engineering*, 1(3/4).
31. Singh, M. P. (1999). An ontology for commitments in multiagent systems. *Artificial Intelligence and Law*, 7(1).